

Thinking about Baseball

Chase Mathis

Introduction

When looking at baseball and predicting the results of baseball games, many factors are at play:

- Who is pitching?
- Is the star player hurt?
- Is the baseball team on a hot streak or a cold run?

This is only a few out of many different factors. Sports betting and casino's now don't create lines based off factors, but are lazy and will depend on the [wisdom of the masses](#). They engineer lines to get equal attraction on both sides of the spread and make money off of the margin.

Importing Packages

```
library(tidyverse) # data wrangling
library(tidymodels) # modeling
library(baseballr) # api for baseball data
library(knitr) # output
library(patchwork)
library(corr) # correlation

ggplot2::theme_set(ggplot2::theme_minimal(base_size = 15))
```

Data

I will retrieve the data from the [baseballr](#) package, where we can scrape data from [baseball-reference.com](#). Instead of looking at one team and predicting off of their statistics (may be done later with a more machine learning approach), I will attempt to make a more general model of which factors make the most predictive model for any baseball team.

I randomly sampled six numbers that correspond to the final standings of last year's season. These will be the teams where I get the data.

```
standings <- read_csv("data/standings.csv") %>%  
  select(Rk, Tm)
```

Rows: 31 Columns: 24

-- Column specification -----

Delimiter: ","

chr (14): Tm, pythWL, vEast, vCent, vWest, Inter, Home, Road, ExInn, 1Run, v...

dbl (10): Rk, W, L, W-L%, R, RA, Rdiff, SOS, SRS, Luck

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
set.seed(123)  
  
randnums <- sample(1:30, 6, replace=FALSE)  
  
standings %>%  
  filter(Rk %in% randnums) %>%  
  select(Tm) %>%  
  kable(col.names = "Teams Randomly Selected")
```

Teams Randomly Selected

Tampa Bay Rays
St. Louis Cardinals
Cincinnati Reds
Philadelphia Phillies
Detroit Tigers
New York Mets

```

teams <- standings %>%
  filter(Rk %in% randnums) %>%
  select(Tm) %>%
  pull()

astros0 <- team_results_bref("HOU", 2021) #baseball reference
mets0 <- team_results_bref("NYM", 2021) #baseball reference
marlins0 <- team_results_bref("MIA", 2021) #baseball reference
nationals0 <- team_results_bref("WSN", 2021) #baseball reference
pirates0 <- team_results_bref("PIT", 2021) #baseball reference
diamond0 <- team_results_bref("ARI", 2021) #baseball reference

baseball<-bind_rows(astros0,mets0,marlins0,nationals0,pirates0,diamond0)

```

Our final data frame:

```
baseball
```

```
-- MLB Team Results data from baseball-reference.com ----- baseballr 1.2.0 --
```

```
i Data updated: 2022-05-16 09:50:08 PDT
```

```
# A tibble: 972 x 22
```

	Gm	Date	Tm	H_A	Opp	Result	R	RA	Inn	Record	Rank	GB
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<chr>
1	1	Thursday~	HOU	A	OAK	W	8	1	" "	1-0	1	Tied
2	2	Friday, ~	HOU	A	OAK	W	9	5	" "	2-0	1	up 1~
3	3	Saturday~	HOU	A	OAK	W	9	1	" "	3-0	1	up 1~
4	4	Sunday, ~	HOU	A	OAK	W	9	2	" "	4-0	1	up 1~
5	5	Monday, ~	HOU	A	LAA	L	6	7	" "	4-1	1	Tied
6	6	Tuesday,~	HOU	A	LAA	W	4	2	" "	5-1	1	up 1~
7	7	Thursday~	HOU	H	OAK	W	6	2	" "	6-1	1	up 1~
8	8	Friday, ~	HOU	H	OAK	L	2	6	" "	6-2	1	Tied
9	9	Saturday~	HOU	H	OAK	L	3	7	" "	6-3	1	Tied
10	10	Monday, ~	HOU	H	DET	L	2	6	" "	6-4	2	1.0

```

# ... with 962 more rows, and 10 more variables: Win <chr>, Loss <chr>,
#   Save <chr>, Time <chr>, `D/N` <chr>, Attendance <dbl>, cLI <dbl>,
#   Streak <dbl>, Orig_Scheduled <chr>, Year <dbl>

```

Given that we don't care how each team individually is predicted, I will remove the `Tm` variable. In addition, the `Orig_Scheduled` has no real values so we can remove that. The `Year` variable has no variance (all 2021), so we can remove that.

In addition, I will input a 9 for blank innings and make it a double; create a variable called `margin` that is the margin of victory or defeat; lastly I will remove some levels of the win or loss variable.

Baseball-reference denotes a walk off win differently than a regular win. To make things simpler, I will remove this distinction


```
baseball <- baseball %>%
  select(-c(Tm,Orig_Scheduled,Year)) %>%
  mutate(Inn = as.double(if_else(str_detect(Inn, ""), Inn, "9")),
         margin = R-RA,
         Result = if_else(str_detect(Result,"L"), "L", "W"),
         Won = if_else(Result == "W", 1, 0),
         Day = if_else(`D/N` == "D", 1,0),
         across(where(is.character),as_factor))
```

Exploratory Data Analysis

Distribution Analysis

```
baseball %>%  
  ggplot(aes(x = Result)) +  
  geom_histogram(stat = 'count') +  
  labs(  
    x = "Win or Loss",  
    y = "Count",  
    title = "Distribution of our Response Variable"  
  )
```

Warning: Ignoring unknown parameters: binwidth, bins, pad




We randomly selected six team that have lost more than won. This can be an issue with logistic regression, but there are ways to fix it as seen in the recipe step.

Correlation Analysis

```
correlation_table <- baseball %>%
  correlate() %>%
  select(term, Won) %>%
  drop_na()

correlation_table %>%
  ggplot(aes(x = fct_reorder(term, Won), y = Won)) +
  geom_col() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  ylim(-1,1) +
  labs(
    x = "Potential Predictor",
    y = "Correlation Coefficient",
    title = "Predictors' Correlation with the a Win"
  )
```



```
# change won to factor
baseball <- baseball %>%
  mutate(Won = as_factor(Won))
```

The correlation graph mostly makes sense. Runs are highly positively correlated with a win, while runs against are highly negatively correlated with the winning a baseball game.

Modeling

Training/testing split

```
set.seed(123)
baseball_split <- initial_split(baseball)
baseball_training <- training(baseball_split)
baseball_testing <- testing(baseball_split)
```

Specify a Model

```
baseball_spec <- logistic_reg() %>%
  set_engine("glm")
```

Recipe 1

```
bball_rec_1 <- recipe(Won ~ Rank + Date, data = baseball_training) %>%
  update_role(Date, new_role = "id variable") %>%
  step_dummy(all_nominal_predictors()) %>% # dummy coding
  step_zv(all_predictors()) # remove zero variance variables

bball_rec_1
```

Recipe

Inputs:

	role	#variables
id variable		1

outcome	1
predictor	1

Operations:

Dummy variables from `all_nominal_predictors()`
 Zero variance filter on `all_predictors()`

Recipe 2

```
bball_rec_2 <- recipe(Won ~ Date + Rank + cLI, data = baseball_training) %>%
  update_role(Date, new_role = "id variable") %>%
  step_dummy(all_nominal_predictors()) %>% # dummy coding
  step_zv(all_predictors()) # remove zero variance variables

bball_rec_2
```

Recipe

Inputs:

	role	#variables
id variable		1
outcome		1
predictor		2

Operations:

Dummy variables from `all_nominal_predictors()`
 Zero variance filter on `all_predictors()`

Recipe 3

```
bball_rec_3 <- recipe(Won ~ Date + Rank + cLI + Day, data = baseball_training) %>%
  update_role(Date, new_role = "id variable") %>%
  step_dummy(all_nominal_predictors()) %>% # dummy coding
  step_zv(all_predictors()) # remove zero variance variables
```


Create Workflows

```
# workflow brings together recipe and model
bball_wflow1 <- workflow() %>%
  add_model(baseball_spec) %>%
  add_recipe(bball_rec_1)

# workflow brings together recipe and model
bball_wflow2 <- workflow() %>%
  add_model(baseball_spec) %>%
  add_recipe(bball_rec_2)

bball_wflow3 <- workflow() %>%
  add_model(baseball_spec) %>%
  add_recipe(bball_rec_3)
```

Perform Cross Validation

```
set.seed(345) # set seed

baseball_folds <- vfold_cv(baseball_training, v = 5) # split up data into 5 partitions

bball_fit_cv1 <- bball_wflow1 %>%
  fit_resamples(baseball_folds,
    control = control_resamples(save_pred = TRUE)) # fit all the folds using rec

bball_fit_cv2 <- bball_wflow2 %>%
  fit_resamples(baseball_folds,
    control = control_resamples(save_pred = TRUE)) # fit all the folds using rec

bball_fit_cv3 <- bball_wflow3 %>%
  fit_resamples(baseball_folds,
    control = control_resamples(save_pred = TRUE)) # fit all the folds using rec

collect_metrics(bball_fit_cv1) %>%
  select(.metric, mean) %>%
  kable(digits = 3)
```

.metric	mean
accuracy	0.595
roc_auc	0.596

```
collect_metrics(bball_fit_cv2) %>%
  select(.metric,mean) %>%
  kable(digits = 3)
```

.metric	mean
accuracy	0.586
roc_auc	0.591

```
collect_metrics(bball_fit_cv3) %>%
  select(.metric,mean) %>%
  kable(digits = 3)
```

.metric	mean
accuracy	0.583
roc_auc	0.589

Visualizing the ROC Curve

Recipe 1 through 3 shown below.

```
p1 <- bball_fit_cv1 %>%
  collect_predictions() %>%
  group_by(id) %>%
  roc_curve(
    truth = Won,
    .pred_0
  ) %>%
  autoplot()

p2 <- bball_fit_cv2 %>%
  collect_predictions() %>%
  group_by(id) %>%
  roc_curve(
```


```

    truth = Won,
    .pred_0
  ) %>%
  autoplot()


p3 <- bball_fit_cv3 %>%
  collect_predictions() %>%
  group_by(id) %>%
  roc_curve(
    truth = Won,
    .pred_0
  ) %>%
  autoplot()

```


p1



p2



p3



Model Selection

Clearly Model 1 is the best model and also the most parsimonious. It has the most consistent ROC Curve, and the highest metrics for accuracy.

Apply Model to testing Data

```
final_fit <- last_fit(
  bball_wflow1,
  split = baseball_split
)

collect_predictions(final_fit) %>%
  conf_mat(Won, .pred_class)
```

Truth

Prediction	0	1
0	116	79
1	14	34

```
collect_metrics(final_fit) %>%
  kable(digits = 3)
```

.metric	.estimator	.estimate	.config
accuracy	binary	0.617	Preprocessor1_Model1
roc_auc	binary	0.647	Preprocessor1_Model1

By simply looking at the baseball's team rank in their respective division, we can achieve a 61.7% accuracy in predicting whether the baseball team will win or lose.